

## Projects descriptions and indications

You are asked to form groups of 2/3 students max. One student per group has to communicate by email ([dario.colazzo@polytechnique.org](mailto:dario.colazzo@polytechnique.org)) the group composition, **by January 30**.

Each group has to chose one of the following problems, design algorithms for the solution, and write a report including:

- a description of the adopted solution
- designed algorithms plus related comments/description; comments to main fragments of code
- experimental analysis, concerning in particular scalability, except for the last project on recommendation systems.
- comments about the experimental analysis outlining weak and strong points of the algorithms.
- an appendix including the code.

A pdf version of the report has to be sent via email ([dario.colazzo@polytechnique.org](mailto:dario.colazzo@polytechnique.org)) before march 25.

For any question regarding the first 4 projects feel free to ask by email if needed ([dario.colazzo@polytechnique.org](mailto:dario.colazzo@polytechnique.org)), while for the project on recommendation system please ask questions to Maryan Morel ([morel.maryan@gmail.com](mailto:morel.maryan@gmail.com)).

**Matrix multiplication.** Consider the problem of multiplying two big matrixes  $A_{n,m}$  and  $B_{m,n}$ . Design MapReduce algorithms for this problem, knowing that up to two jobs may be needed. Consider the simple textual representation of a matrix where the value  $v$  for  $A_{i,j}$  is represented by a text line 'i j v'. Only non-zero values will have a line in the textual file representing a matrix. Provide at least two non-trivial different algorithms, and implement them in both Python with Hadoop streaming and in Spark. Perform experimental evaluation by considering 5 couples of matrixes of increasing sizes, by doubling the sizes from a pair of matrix to the subsequent pair. You can use Spark to generate matrixes.

**TF-IDF.** Consider the problem of calculating the TF-IDF score for each word in a set of documents for each document. Provide a MapReduce algorithm to calculate TF-IDF scores given an input set of documents. Provide both a MapReduce Python Hadoop streaming and a Spark implementation. Perform experimental analysis in order to compare performances of the two implementation. To test scalability consider 5 document sets of increasing sizes. For instance, size can double from a set to another.

The set of input textual document can be either dowloaded from the Web or generated by a Python script for instance. To speed up document generation MapReduce can be used, for instance by requiring that  $n$  Reducers generate a certain amount of documents each. Words of documents can be randomly picked from an input fixed vocabulary. Also, you can use Spark to make document generation easier.

**Words co-occurrence matrix.** Consider the the problem of building a word co-occurrence matrice from large corpora of textual documents. Formally, the co-occurrence matrix is a square  $n \times n$  matrix where  $n$  is the number of unique words in all the documents (i.e., the vocabulary size). A cell  $m_{ij}$  contains the number of times word  $w_i$  co-occurs with word  $w_j$  within a specific context—a natural unit such as a sentence, paragraph, or a document; for simplicity consider that the context is a document, and consider an input collection of relatively small documents, that you can generate starting from a fixed vocabulary. MapReduce can be used to generate the documents. Note that the upper and lower triangles of the matrix are identical since co-occurrence is a symmetric relation.

Define MapReduce algorithms for this problem and implement them both in Python-Hadoop streaming and Spark, and compare them on five sets of increasing sizes. For instance, size can double from a set to another.

Also, you can use Spark to make document generation easier.

**Single-source shortest path problem.\*** The task is to find shortest paths from a source node to all other nodes in the graph. This problem is solved by the Dijkstra's algorithm, which is sequential. The project has a double purpose. First get familiar with Dijkstra's algorithm, then devise a MapReduce version of the algorithm. As you will realise, the process is actually iterative, so the identified MapReduce job must be iterated a certain number of times.

Provide both a Python-Hadoop streaming and Spark implementation of the algorithm, and test it on the simple graph data provided in classes.

Optional: perform scalability experiments as for previous projects. A single comparison on a reasonable big graph would be sufficient

### **Simple recommender system based on latent classes - Probabilistic Semantic Latent Indexing (PLSI)**

Let us consider a dyadic dataset, composed of  $N$  users  $U = \{u_1, \dots, u_N\}$  and  $M$  items  $I = \{i_1, \dots, i_M\}$ . For each (user, item) pair, you know if a user interacted with or liked an item. In this project, you will implement a basic recommendation algorithm based on Hoffman (1999) work, to recommend items to users.

Hoffman (1999) describes a simple probabilistic algorithm based on latent classes. Latent classes are unobserved classes, clustering both users and items. The authors assume a multinomial distribution of users given latent classes, and of items given classes. Users and item can then belong to several latent classes. Latent classes are expected to gather similar items and users, for example, there can be a class gathering sci-fi fans and Star Trek movies.

Parameter estimation is done using the EM algorithm, which is a standard procedure when working with latent parameters.

This algorithm was one of the building bricks of Google News recommender system. Das (2007) describe Google's MapReduce implementation of the basic PLSI algorithm (see section 4.2).

You are asked to implement in *Spark* the PLS algorithm, based on Das (2007) MapReduce formulation, and to apply it on the well known Movielens dataset.

Note that the Movielens dataset contains ratings ranging from 1 to 5. You can reduce this information to seen/not seen in order to use the basic version of PLSI. If you are very motivated, you can try to integrate preference values to PLSI, as described in Hoffman (1999). You can also compare your results to Spark LDA implementation and discuss the use of LDA versus PLSI.

---

[1] Hofmann, Thomas, and Jan Puzicha. "Latent class models for collaborative filtering." IJCAI. Vol. 99. No. 1999. 1999. – <http://people.eecs.berkeley.edu/~russell/classes/cs294/f05/papers/hofmann+puzicha-1999.pdf>

[2] Das, A. S., Datar, M., Garg, A., & Rajaram, S. (2007, May). Google news personalization: scalable online collaborative filtering. In Proceedings of the 16th international conference on World Wide Web (pp. 271-280). ACM. – <https://pdfs.semanticscholar.org/fb42/07376177f18a6cf58b53ecd231fb3395ca33.pdf>

[3] <https://grouplens.org/datasets/movielens/>

[4] <https://spark.apache.org/docs/2.2.0/api/python/pyspark.ml.html#pyspark.ml.clustering.LDA>

IMPORTANT, for all projects: you are allowed to consult available documentation on the Web in order to learn required algorithms, in particular for the fourth project.

A suggested text is *Data-Intensive Text Processing with MapReduce*, by Jimmy Lin and Chris Dyer, freely available on the web, Just google the title plus 'pdf'.

Another important aspect is that in each project (except the last one) you are required to take particular care of scalability issues and in particular on implementing techniques for local aggregation, and for minimising the amount of information sent to the shuffle phase.